TableNet with a wheel history encoder

one network for picks and wheels · `train_tablenet_joint.py whist_layers=2`

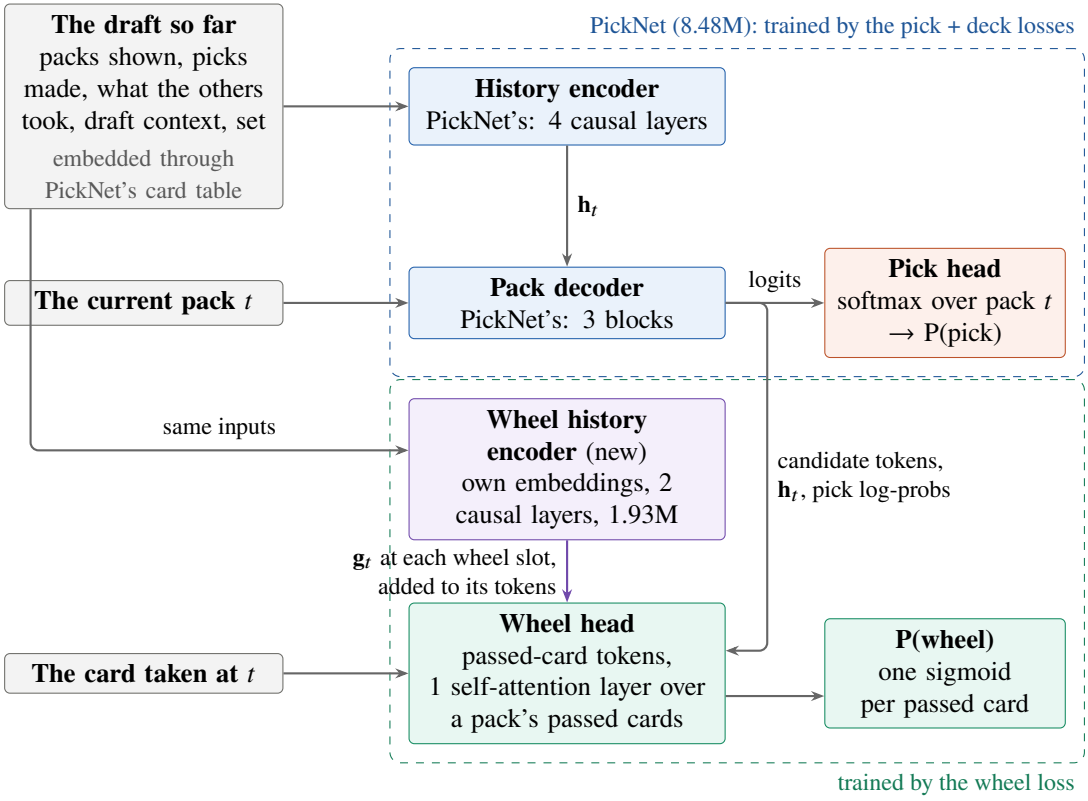

Why the new box. With one shared history encoder, the pick and wheel losses pull it in opposite directions, and the wheel task's extra accuracy came only from reshaping it, at a cost to picks. The wheel history encoder gives the wheel head its own summary of the draft. It reads the same inputs as PickNet's, through PickNet's card table, so PickNet's history can stay tuned for picks. Its output projection starts at zero, so training begins exactly at the step-1 wheel head.

Who trains what. PickNet always trains on the pick and deck losses. The wheel loss trains the wheel head and the wheel history encoder. How far the wheel loss also reaches into PickNet is set by `wheel_block`: `all` keeps it out of PickNet entirely (run A); `hist` keeps it out of PickNet's history encoder only, so it can still adjust the card table and pack decoder (run B).